

# Mixed Precision

Christoph Höppke, Daniel Tomaschewski

TU Dortmund

Version: 18. Mai 2016

# Content

---

## 1 Mixedprecision Definition and Goals

- Definition
- Precision
- Data Error and Truncation
- Floting Point Operations. A deeper analysis

## 2 History of Mixed Precision in the context of GPGPU calculations

- History of GPGPU
- Why Mixed Precision is difficult
- Unconventional computation Hardware

## 3 Example Calculation

- Testresults

# Content

---

## 1 Mixedprecision Definition and Goals

- Definition
- Precision
- Data Error and Truncation
- Floting Point Operations. A deeper analysis

## 2 History of Mixed Precision in the context of GPGPU calculations

- History of GPGPU
- Why Mixed Precision is difficult
- Unconventional computation Hardware

## 3 Example Calculation

- Testresults

# Definition

---

## Definition:

An Algorithm that uses different precisions in its computation

## Goal:

Obtain the **same accuracy** by using high precision  
but **better performance** by utilizing low precision computations

## Performance Gains for bandwidth bound algorithms

- 64 bit = 1 double = 2 floats = 4 halves
- More variables per **bandwidth** and variables per **storage**
- Applies to all memory levels: network, disc, main, device, local, register

## Performance Gains for computation bound algorithms

- 1 double addition  $\approx$  2 float additions  $\approx$  4 half additions (linear)
- 1 double multip.  $\approx$  4 float multip.  $\approx$  16 half multip. (quadratic)
- $\rightarrow$  up to 16 times better computational efficiency

# Definition

---

## Definition:

An Algorithm that uses different precisions in its computation

## Goal:

Obtain the **same accuracy** by using high precision  
but **better performance** by utilizing low precision computations

## Performance Gains for bandwidth bound algorithms

- 64 bit = 1 double = 2 floats = 4 halves
- More variables per **bandwidth** and variables per **storage**
- Applies to all memory levels: network, disc, main, device, local, register

## Performance Gains for computation bound algorithms

- 1 double addition  $\approx$  2 float additions  $\approx$  4 half additions (linear)
- 1 double multip.  $\approx$  4 float multip.  $\approx$  16 half multip. (quadratic)
- $\rightarrow$  up to 16 times better computational efficiency

# Definition

---

## Definition:

An Algorithm that uses different precisions in its computation

## Goal:

Obtain the **same accuracy** by using high precision  
but **better performance** by utilizing low precision computations

## Performance Gains for bandwidth bound algorithms

- 64 bit = 1 double = 2 floats = 4 halves
- More variables per **bandwidth** and variables per **storage**
- Applies to all memory levels: network, disc, main, device, local, register

## Performance Gains for computation bound algorithms

- 1 double addition  $\approx$  2 float additions  $\approx$  4 half additions (linear)
- 1 double multip.  $\approx$  4 float multip.  $\approx$  16 half multip. (quadratic)
- $\rightarrow$  up to 16 times better computational efficiency

# Definition

---

## Definition:

An Algorithm that uses different precisions in its computation

## Goal:

Obtain the **same accuracy** by using high precision  
but **better performance** by utilizing low precision computations

## Performance Gains for bandwidth bound algorithms

- 64 bit = 1 double = 2 floats = 4 halves
- More variables per **bandwidth** and variables per **storage**
- Applies to all memory levels: network, disc, main, device, local, register

## Performance Gains for computation bound algorithms

- 1 double addition  $\approx$  2 float additions  $\approx$  4 half additions (linear)
- 1 double multip.  $\approx$  4 float multip.  $\approx$  16 half multip. (quadratic)
- $\rightarrow$  up to 16 times better computational efficiency

# A brief history overview

---

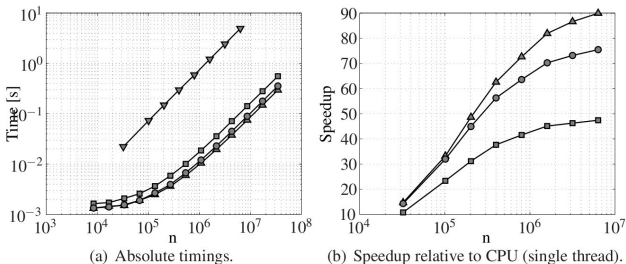
## Trends:

- Memoryclock speeds are increasing
  - GTX 980 Ti Memoryclock speed: 2x 1753 MHz
  - GTX 1080 Memoryclock speeds: 4x 2500 MHz (+185%)
- Alternative computing arcetectures
  - APU's
  - SOC's
  - → Sharing memory between CPU and GPU is becoming easier and more common
- NVIDIA pushing the use of half precision.
  - half and half2 were announced as important new features in the CUDA Toolkit version 7.5



# A brief history overview

## Past achievements:



**Fig. 3** Scalability tests and performance comparisons on Tesla C2050 in single precision (—▲—), double precision (—■—), mixed precision (—●—), and CPU (single thread) code (—▼—). Sixth order spatial discretization employed. The iterative defect correction method has been left-preconditioned with a Gauss-Seidel V-cycle multigrid strategy on each architecture.

# Roundoff and Cancellation

## Definition Machine Precision

The smallest positive number  $\epsilon$  for which a floating point calculation evaluates the expression  $1 + \epsilon > 1$  to be true.

### Examples:

- $\epsilon_{double} \approx 2.220446049250313 \cdot 10^{-16}$
- $\epsilon_{float} \approx 1.1920929 \cdot 10^{-7}$
- $\epsilon_{half} \approx 9.765625 \cdot 10^{-4}$
- So **more precision** is usually **better**

## Cancellation

additive roundoff	$a = 1 + 0.00000004 = 1.00000004$	$=_{fl} 1$
multiplicative roundoff	$b = 1.0002 \cdot 0.9998 = 0.99999996$	$=_{fl} 1$
<b>cancellation</b>	$c \in \{a, b\}$	$\pm 4 = (c - 1) \cdot 10^8 =_{fl} 0$
order of operations	$1 + 0.00000004 - 1 =_{fl} 0$	
	$1 - 1 + 0.00000004 =_{fl} 0.00000004$	

# Roundoff and Cancellation

## Definition Machine Precision

The smallest positive number  $\epsilon$  for which a floating point calculation evaluates the expression  $1 + \epsilon > 1$  to be true.

### Examples:

- $\epsilon_{double} \approx 2.220446049250313 \cdot 10^{-16}$
- $\epsilon_{float} \approx 1.1920929 \cdot 10^{-7}$
- $\epsilon_{half} \approx 9.765625 \cdot 10^{-4}$
- So **more precision** is usually **better**

## Cancellation

additive roundoff       $a = 1 + 0.00000004 = 1.00000004 =_{fl} 1$

multiplicative roundoff       $b = 1.0002 \cdot 0.9998 = 0.99999996 =_{fl} 1$

**cancellation**       $c \in \{a, b\} \quad \pm 4 = (c - 1) \cdot 10^8 =_{fl} 0$

order of operations       $1 + 0.00000004 - 1 =_{fl} 0$   
                                  $1 - 1 + 0.00000004 =_{fl} 0.00000004$

# Roundoff and Cancellation

## Definition Machine Precision

The smallest positive number  $\epsilon$  for which a floating point calculation evaluates the expression  $1 + \epsilon > 1$  to be true.

### Examples:

- $\epsilon_{double} \approx 2.220446049250313 \cdot 10^{-16}$
- $\epsilon_{float} \approx 1.1920929 \cdot 10^{-7}$
- $\epsilon_{half} \approx 9.765625 \cdot 10^{-4}$
- So **more precision** is usually **better**

### Cancellation

additive roundoff	$a = 1 + 0.00000004 = 1.00000004$	$=_{fl} 1$
multiplicative roundoff	$b = 1.0002 \cdot 0.9998 = 0.99999996$	$=_{fl} 1$
<b>cancellation</b>	$c \in \{a, b\}$	$\pm 4 = (c - 1) \cdot 10^8 =_{fl} 0$
order of operations	$1 + 0.00000004 - 1 =_{fl} 0$	
	$1 - 1 + 0.00000004 =_{fl} 0.00000004$	

# Computational Precision vs Accuracy of Result

---

Instructive Example [S.M. Rump, 1988]

$$f(x, y) = (333.75 - x^2)y^6 + x^2(11x^2y^2 - 121y^4 - 2) + 5.5y^8 + 0.5x/y$$
$$x_0 = 77617, y_0 = 33096$$

float s23e8

1.1726

double s52e11

1.17260394005318

The correct result is:

-0.82739605994682136814116509547981629

# Computational Precision vs Accuracy of Result

---

Instructive Example [S.M. Rump, 1988]

$$f(x, y) = (333.75 - x^2)y^6 + x^2(11x^2y^2 - 121y^4 - 2) + 5.5y^8 + 0.5x/y$$
$$x_0 = 77617, y_0 = 33096$$

float s23e8

1.1726

double s52e11

1.17260394005318

**The correct result is:**

-0.82739605994682136814116509547981629

# DataError and Truncation

---

- Data error occurs when the **exact value** has to be **truncated** for storage in the binary format
  - $\pi$  ,  $\sqrt{2}$ ,  $\sin(2)$ ,  $e^2$ ,  $1/3$
  - every rational number with a denominator that has a prime factor other than 2
- How can float be better than double?
  - There is **no data error** in the operands
  - The errors can **cancel out** themselves favorably

# Floating Point Operations. A deeper analysis

**Number representation** → almost all numbers have to be truncated

half s10e5 a	1 bit sign $s_a$	10 bit mantissa $m_a$	5 bit exp. $e_a$
float s23e8 b	1 bit sign $s_b$	23 bit mantissa $m_b$	8 bit exp. $e_b$
double s52e11 c	1 bit sign $s_c$	52 bit mantissa $m_c$	11 bit exp. $e_c$

Multiplication		
Precision	Exactformat	main error:
half(a) · half(b)	s20e6	truncate mantissa from 20 to 10 bit
float(a) · float(b)	s46e9	truncate mantissa from 46 to 23 bit
double(a) · double(b)	s104e12	truncate mantissa from 104 to 52 bit
Addition		
Precision	Exactformat	main error:
half(a) · half(b)	s41e5	truncate mantissa from 41 to 10 bit
float(a) · float(b)	s278e8	truncate mantissa from 278 to 23 bit
double(a) · double(b)	s2099e11	truncate mantissa from 2099 to 52 bit



# Content

---

## 1 Mixedprecision Definition and Goals

- Definition
- Precision
- Data Error and Truncation
- Floting Point Operations. A deeper analysis

## 2 History of Mixed Precision in the context of GPGPU calculations

- History of GPGPU
- Why Mixed Precision is difficult
- Unconventional computation Hardware

## 3 Example Calculation

- Testresults

# History of GPGPU

---

## **1999: NVIDIA GeForce 256**

- Term GPU was coined during the launch
- GPGPU calculations were only achievable by "Hacking the GPU"(very cumbersome and error-prone)

## **2003-2007 First wave of GPU computing**

- Floating point support
- Performance improvements
- Geforce 8800GT Release (29.October 2007)  
first CUDA enabled Consumer GPU
- Double precision was not available on the GPU
- Mixed Precision was the ONLY way to utilize GPU horsepower without precision compromises
- Mixed Precision lead to a speedup of factor 3-5 in FEM applications

# Why Mixed Precision is difficult

---

- Data has to be transferred to the GPU over the relatively **slow PCIe** interface
- A lot of time is wasted while waiting for data

# Unconventional computation Hardware

---

## **NVIDIA Jetson TK1 Bord**

- CPU and GPU share the same memory and use the same memory interface
- Copying data from CPU to GPU and vice versa is much faster

# Content

---

## 1 Mixedprecision Definition and Goals

- Definition
- Precision
- Data Error and Truncation
- Floting Point Operations. A deeper analysis

## 2 History of Mixed Precision in the context of GPGPU calculations

- History of GPGPU
- Why Mixed Precision is difficult
- Unconventional computation Hardware

## 3 Example Calculation

- Testresults

# Example Calculation

---

## Problem

Solve the Poisson problem  $-\Delta u = 1$ ,  $x \in \Omega$  with dirichlet boundry conditions  $u \equiv 0$ ,  $x \in \partial\Omega$  using conforming quadrilateral elements for the finite-element discretization of the unit square  $\Omega = (-1, 1)^2$

- Method 1: Using a CG solver
- Method 2 : Using a more sophisticated MG solver

# Testresults pending

---

Imagine a nice Plot